

Lecture 12: Interval Trees, Segment Trees

Prof. Nodari Sitchinava

Scribe: Mojtaba Abolfazli, Muzamil Yahia

1 Priority search tree: cont

This section is a continuation of the last session on Priority Search Trees (PSTs) where we are interested to dive deeper into the analysis of PSTs. Before going into details, we need to recall PST with P_{ymax} as the point with the largest y -coordinate and x_{med} as the median x -coordinate among points in database S . A PST can be built up as shown in Figure 1 where S_{left} and S_{right} are obtained as follows:

$$S_{left} = \{P \in S \setminus \{P_{ymax}\} \mid P.x \leq x_{med}\}$$

$$S_{right} = \{P \in S \setminus \{P_{ymax}\} \mid P.x > x_{med}\}$$




Figure 1: Building priority search tree.

To report all nodes in a subtree rooted at v , we can use the following pseudocode:

Algorithm 1 Report a subtree in PST

```

1: function REPORT( $v$ )
2:   if  $v \neq NULL$  then
3:     OUTPUT( $v$ )
4:     REPORT( $v.left$ )
5:     REPORT( $v.right$ )

```

1.1 Range Query

Let's consider a range query in one dimensional space where we're asked to report all points greater than or equal to a real number q_y . In this case, we can use the following pseudocode to report all nodes above q_y in a subtree rooted at v .

Algorithm 2 Report a subtree above q_y in PST

```
1: function REPORT( $v, q_y$ )
2:   if  $v \neq NULL$  then
3:     if  $v.y \geq q_y$  then
4:       OUTPUT( $v$ )
5:       REPORT( $v.left, q_y$ )
6:       REPORT( $v.right, q_y$ )
```

The above pseudocode reports all the points above q_y because any path from root v to any node is in a non-decreasing order of y -coordinate (otherwise the property of PST doesn't hold). Likewise, the subtree of nodes needs to be reported is either connected and contains the root or it is empty.

The performance of PSTs is summarized in the following theorem.

Theorem 1. *The total running time to report all points in a query range $[q_y, +\infty]$ in a PST rooted at v is:*

$$T(v) = c \cdot k_v + c - 1$$

where k_v is the number of reported points and $c \geq 1$.

Proof. We prove the theorem by using the induction.

Induction hypothesis: For all descendant nodes u of v :

$$T(u) = c \cdot k_u + c - 1$$

where k_u is the total number of nodes above q_y in a PST rooted at u .

For the base case where there is no point in the query above q_y , the theorem is true. Now we show that it's also true for the PST rooted at v .

$$\begin{aligned} T(v) &= 1 + T(v.left) + T(v.right) \\ &= 1 + c \cdot k_{v.left} + c - 1 + c \cdot k_{v.right} + c - 1 \\ &= c - 1 + c(1 + k_{v.left} + k_{v.right}) \\ &= c - 1 + c \cdot k_v \end{aligned}$$

□

1.2 3-sided Range Query

Now, let's look at 2D range query where we aim to find nodes within a 3-sided range (i.e. $[x_L, x_R] \times [y, +\infty]$) as shown in Figure 2.

We follow the subsequent steps to report all nodes corresponding to the specified range:

- Find splitting vertex v_{split} .
- Follow left path down to x_L (left boundary) and report all points in the subtrees hanging from the right.




Figure 2: 3-sided Range Query.

- Follow right path down to x_R (right boundary) and report all points in the subtrees hanging from the left.

The pseudocode shown in Algorithm 3 describes the query algorithm:

We start to find the splitting point with x_L and x_R as shown in Figure 3. All shaded subtrees are those whose x -coordinate lies in the specified range. Therefore, we can search those subtrees based on just y -coordinate as given in Algorithm 2.




Figure 3: Querying a PST.

As discussed in the previous session, the running time for 2D query is of the order $T(v) = O(\log n + k)$ where k denote the number of reported points. The PST for n points uses $O(n \log n)$ memory space to store the data and takes $O(n \log n)$ to build.

2 Interval trees for orthogonal stabbing queries

In certain geometric applications like planar graphs, we might be interested in reporting the set of line segments L associated with a window query $w = [x_0, x_1] \times [y_0, y_1]$. This set L contains all line segments with at least one end point inside w , and those segments that span the window w without having an end point inside w .

Algorithm 3 3-sided Range Query

```
1: function 3-SIDED QUERY( $v, [x_L, x_R], y$ )
2:   while ( $v \neq NULL$ ) and ( $v.x_{med} \leq x_L$  or  $x_R < v.x_{med}$ ) do
3:     if  $v \in [x_L, x_R] \times [y, +\infty]$  then
4:       OUTPUT( $v$ )
5:     if  $v.x_{med} \leq x_L$  then
6:        $v = v.right$ 
7:     else
8:        $v = v.left$ 
9:      $v_{split} = v$ 
10:     $v = v.left$ 
11:    if  $v_{split} \in [x_L, x_R] \times [y, +\infty]$  then
12:      OUTPUT( $v_{split}$ )
13:  while  $v \neq NULL$  do
14:    if  $v \in [x_L, x_R] \times [y, +\infty]$  then
15:      OUTPUT( $v$ )
16:    if  $x_L \leq v.x_{med}$  then
17:      REPORT( $v.right, y$ )
18:       $v = v.left$ 
19:    else
20:       $v = v.right$ 
21:       $v = v_{split}.right$ 
22:  while  $v \neq NULL$  do
23:    if  $v \in [x_L, x_R] \times [y, +\infty]$  then
24:      OUTPUT( $v$ )
25:    if  $v.x_{med} \leq x_R$  then
26:      REPORT( $v.left, y$ )
27:       $v = v.right$ 
28:    else
29:       $v = v.left$ 
```

We can report all line segments with at least one endpoint inside w using algorithms and techniques from Section 1. We can use 2 instances of 3-sided query i.e. $q_0 = [x_0, x_1] \times [y_0, +\infty]$ and $q_1 = [x_0, x_1] \times [-\infty, y_1]$, and then report the intersections of the two queries in time $O(\log n + k_0) + O(\log n + k_1)$. An algorithm to find intersections $q_0 \cap q_1$, would be to keep sorted lists of q_0 and q_1 and then find the intersection in linear time $O(k_0 + k_1)$ using modified merge algorithm in mergesort. However, in the worst case scenario, the two queries may report all $n = k_0 + k_1$ endpoints while the intersection may contain only one element. By constructing 2D range trees with storage $O(n \log n)$, and time $O(\log^2 n + k)$ we can report all line segments with ends points in w . Applying fractional cascading can reduce the query time further to $O(\log n + k)$ which is faster than $O(\log n + n)$.

What is left is how to handle efficiently those line segments that span our window query w without any endpoint inside w ? In other words, we are looking for line segments that cross two boundaries of our window query w . If we know how to report all line segments that cross one boundary, say $[x_0, x_1]$, of w , we can easily check if they cross any of the remaining boundaries. We present first




Figure 4: Query $w = [x_0, x_1] \times [y_0, y_1]$.

a solution to a slightly simpler problem that can help us towards reporting those line segments crossing w but with no end point inside w .

Stabbing query: Given a set of intervals $I = \{i_1, \dots, i_n\}$, and a vertical line query $\ell(q = x)$, find all intervals that are crossed by ℓ .

Claim 2. *We can solve stabbing the query in $O(n \log n)$ and $O(n)$ space.*

We apply divide and conquer strategy to solve stabbing query problem. We construct recursively a balanced BST over all end points of intervals in I , i.e. $E_I = \bigcup_{[i_l, i_r] \in I} \{i_l.x, i_r.x\}$. While doing so, we augment this BST with extra data containing intervals from I that contains $v.val$ in their boundaries as shown in Algorithm 5.

Algorithm 4

```

1: function BUILDINT( $I$ )
2:    $v.val = \text{median}(E_I)$  ▷ key stored at node  $v$ 
3:    $I_{mid} \leftarrow \{i \in I \mid i_l.x \leq v.val \leq i_r.x\}$ 
4:    $I_{left} \leftarrow \{i \in I \mid i_r.x < v.val\}$ 
5:    $I_{mid} \leftarrow \{i \in I \mid v.val < i_l.x\}$ 
6:    $v.data \leftarrow I_{mid}$  ▷ augmented data
7:    $v.left \leftarrow \text{BUILDINT}(I_{left})$ 
8:    $v.right \leftarrow \text{BUILDINT}(I_{right})$ 
9:   return  $v$ .
```

Figure 5 shows 8 line segments (intervals) with median being the right end point of interval i_4 . The root of the interval tree divides the set of intervals $I = \{i_1, \dots, i_8\}$ into 3 disjoint sets $I_{mid}, I_{left}, I_{right}$, where the root contains $I_{mid} = \{i_3, i_4, i_5\}$. The left subtree will recursively store $I_{left} = \{i_1, i_2\}$ and the right subtree will recursively store $I_{right} = \{i_6, i_7, i_8\}$.

How do we store I_{mid} ? In order to retrieve all intervals that are stabbed by $q = x$, we can scan the list of intervals stored at node v for intervals that contain x before going to the left or right




Figure 5: $I = I_{left} \cup I_{mid} \cup I_{right}$, and $\text{STAB}(I, \mathbf{x}) = \{i_3, i_4, i_1, i_2\}$.

subtree depending on whether \mathbf{x} is less than or greater than $v.val$. However, it may happen that all n intervals are stored in the root, and checking for intervals containing \mathbf{x} may take $O(n)$ then instead of $O(k)$, the output size!

In order to prevent that, we store I_{mid} in a more efficient way by storing I_{mid} sorted by the x -coordinate of start point in an increasing order, and another copy of I'_{mid} of I_{mid} ordered by the x -coordinate of end point in a decreasing order. To check if $\mathbf{q} = \mathbf{x}$ is contained in some intervals in v , we compare \mathbf{x} with $v.val$ and accordingly scan the list I_{mid} or I'_{mid} till we find an interval that does not contain \mathbf{x} . In Figure 5, at the root of the tree, $I_{mid} = \{i_3 \leq i_4 \leq i_5\}$ and $I'_{mid} = \{i_3 \geq i_5 \geq i_4\}$. The stabbing query would then run over I_{mid} and would scan $i_3 \cap \{\mathbf{x}\} \neq \phi$, $i_4 \cap \{\mathbf{x}\} \neq \phi$, $i_5 \cap \{\mathbf{x}\} = \phi$.

Algorithm 5

```

1: function QUERYINT( $v, \mathbf{x}$ )
2:   if  $\mathbf{x} = v.val$  or  $I_{mid} = \phi$  then return  $I_{mid}$ 
3:   if  $t < v.val$  then
4:     REPORTINT( $I_{mid}, left, \mathbf{x}, \leq$ )  $\triangleright O(k_v + 1)$ 
5:     QUERYINT( $v.left, t$ )
6:   else
7:     REPORTINT( $I'_{mid}, right, \mathbf{x}, \geq$ )  $\triangleright O(k_v + 1)$ 
8:     QUERYINT( $v.right, t$ )
9:   return  $v$ .
```

Algorithm 6

```

1: function REPORTINT( $I, leftRight, \mathbf{x}, comparator$ )
2:   for  $j = 1$  to  $|I|$  do
3:     if  $comparator(I[j].leftRight.x, \mathbf{x}) = true$  then
4:       output  $I[j]$ 
5:     else
6:       break
```

Since, we store only 2 copies of each interval i in the entire tree, interval trees takes $O(n)$ space. To report all intervals containing \mathbf{x} at node v , we need $k_v + 1$ comparisons. The recurrence relation

is then given by

$$\begin{aligned} T(v) &= O(1 + k_v) + \max\{T(v.left), T(v.right)\} \\ &= O(\log n + k) \end{aligned} \quad (\text{because the tree is balanced}).$$

The preprocessing time takes sorting complexity, as we need to build a balanced BST, sort all intervals by start point, and again by end points. We can write the recurrence relation of BUILTINT as follows:

$$\begin{aligned} T(n, h) &= O(n) + T(|I_{left}|, h - 1) + T(|I_{right}|, h - 1) \\ &= O(nh) \\ &= O(n \log n). \end{aligned}$$

Answering the window query $w = [x_0, \times x_1] \times [y_0 \times y_1]$

In order to answer our original window query w , we can modify the data structure of augmented data I_{mid} stored in interval tree to be 1D range tree over y -axis instead of an ordered list over x -axis. Range tree takes $O(|I_{mid}|)$ to store data, $O(\log |I_{mid}| + k_v)$ for querying its data, and $O(|I_{mid}| \log |I_{mid}|)$ for preprocessing. Since we need to query a range tree at each node v visited while querying the interval tree, it takes $\sum_{v \in \text{Path}(\text{QUERYINT}(v, x))} O(\log n + k_v) = O(\log^2 n + k)$ to answering the window query.

3 Segment trees

Segments tree is another data structure that can be used to answer stabbing queries. It is more general than interval trees as it can process any type of non-intersecting non-vertical line segments, but at the price of increased storage $O(n \log n)$. As we did with interval trees we first shows how segment trees answer stabbing queries and then highlight how to answer window queries.

Consider the set of line segments $L = \{\ell_1, \ell_2, \ell_3, \ell_4, \ell_5\}$ in the following diagram:




Figure 6: Set of line segments $L = \{\ell_1, \ell_2, \ell_3, \ell_4, \ell_5\}$ with slabs $S = \{s_{-\infty}, s_1, \dots, s_9, s_{\infty}\}$.

Observe that we cannot always define a total ordering over slanted line segments that preserves the transitivity property. For example, if we consider $\ell_5 \leq \ell_4 \leq \ell_2$ then by transitivity we would have

$\ell_5 \leq \ell_2$. On the other hand, if we consider $\ell_2 \leq \ell'_4 \leq \ell_5$ then we would have $\ell_2 \leq \ell_5$. However, if we divide the plane into slabs shown by the dashed lines, we can define a total ordering within each slab that is transitive.

We can build a balanced BST for slabs based on x -values of end points with leaves

$$S = \{s_{-\infty}, s_1, \dots, s_9, s_{\infty}\},$$

and internal node v being the union of all slabs within the subtree rooted at v . The root is then all x -axis given by the only slab $[-\infty, \infty]$. Next we augment the BST node v with list of line segments ℓ that crosses any slab with its subtree but do not cross the parent slab.

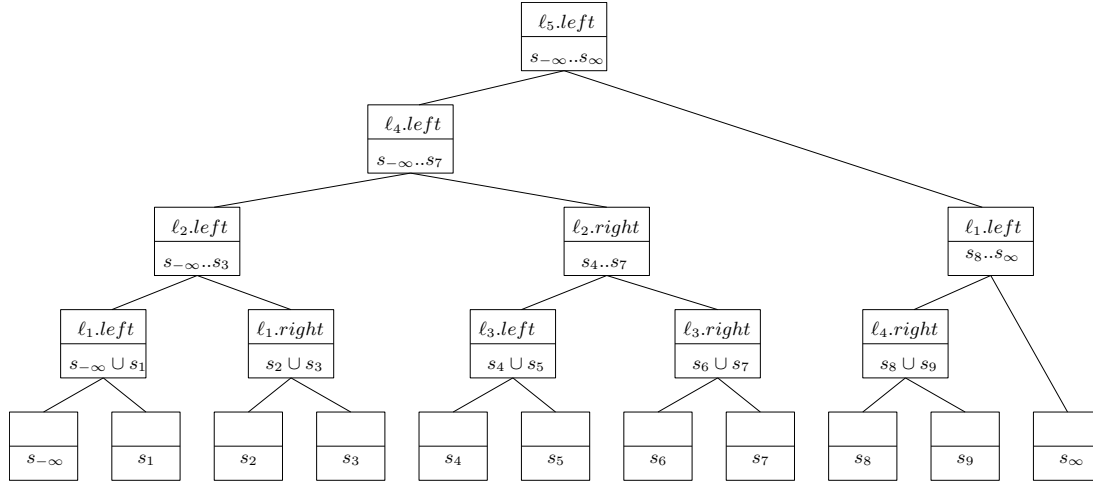


Figure 7: A segment tree for Figure 6 with slabs as leaves, and union of slabs as internal nodes.

Invariant Each line segment ℓ is stored in a node of v such that

$$[v.x_{left}, v.x_{right}] \subseteq [\ell.x_{left}, \ell.x_{right}] \wedge [parent(v).x_{left}, parent(v).x_{right}] \not\subseteq [\ell.x_{left}, \ell.x_{right}].$$

Claim 3. *Each segment will be stored in at most 2 nodes at each level of segment tree.*

The proof of the claim above can be found in chapter 10 in [BCKO08]. Based on the result above, the segment tree would take $O(n \log n)$ space because we store at most $2n$ nodes at each level of the BST, and because the tree is balanced. In order to answer stabbing queries, we locate the slabs that contains the query $q = x$ from root till leaves level (a single branch) and report all line segments containing q . This takes $O(\log n)$ to locate the leaf node, and k to report the output. Hence, it takes $O(\log n + k)$ to answer stabbing query.

Answering the window query $w = [x_0, x_1] \times [y_0, y_1]$

In order to answer our original window query w , we use 1D range tree to store line segments sorted by y -axis. Note that within each slab in segment tree, we have total ordering of line segments and therefore we can build up a BST without affecting the result. The same analysis for interval trees applies here and we would have $O(\log n^2 + k)$ query time for segment trees.

References

- [BCKO08] Mark de Berg, Otfried Cheong, Marc van Kreveld, and Mark Overmars. *Computational Geometry: Algorithms and Applications*. Springer-Verlag TELOS, Santa Clara, CA, USA, 3rd ed. edition, 2008.